

Investigating the effect of Fine Tuning on Bias in Large Language Models

Research Question: How do different fine-tuning approaches affect the performance and bias of Large Language Models in natural language understanding tasks?

Subject: Computer Science

Candidate Code:

Word Count: 3738

Table of Contents

<b>Large Language Models.....</b>	<b>2</b>
<b>Neural Networks.....</b>	<b>2</b>
<b>Transformers.....</b>	<b>3</b>
<b>Bias.....</b>	<b>5</b>
<b>Processing.....</b>	<b>6</b>
<b>Example of Bias.....</b>	<b>7</b>
<b>Prompt Engineering.....</b>	<b>8</b>
<b>Fine-tuning.....</b>	<b>9</b>
<b>Alternative Approaches.....</b>	<b>10</b>
<b>Benefits of RAG.....</b>	<b>10</b>
<b>Experiments.....</b>	<b>12</b>
<b>Future Implications.....</b>	<b>15</b>
<b>Results.....</b>	<b>17</b>
<b>Conclusion.....</b>	<b>21</b>
<b>Bibliography.....</b>	<b>24</b>

## Large Language Models

Large Language Models at a superficial level are systems that can understand and create human-like text. The reason they are called “large” language models can be attributed to the fact that the models are trained on large datasets. LLMs are built on the principle of machine learning: the ability of machines to learn from data and make decisions autonomously. In more complex terms, they utilize a specific type of neural network called a “transformer model.”

## Neural Networks

A neural network is basically an artificial intelligence model that makes decisions similar to a human brain, hence the “neural” in the title. Every neural network has layers of “nodes,” an input layer, one or more hidden layers, and an output layer. All nodes are interconnected and they each have an associated weight and “threshold value.” If the output for any individual node is higher than the threshold value then that node becomes active, sending some information to the next layer of the network. To improve the accuracy of these networks or ensure they identify whatever input you’re giving it you have to feed it training data.

# Transformers

A transformer is a specific type of neural network that alters an input sequence into an output sequence. This is done by analyzing the context and making connections between sequence components. For example, say the user inputs the sequence What color is the sky?, the transformer will compute mathematical relationships between each word and will use that to generate an output. A renowned transformer model today is “Chat-GPT,” which is short for “Generative Pre-training Transformer.” Transformers have different stages in their layers and the first is “input embeddings.” In this stage, the input given by the user is converted to a series of “tokens.” If the user typed in a sentence each token would be a specific word in that sentence. The embeddings inherent in the transformer will take these individual tokens and convert them into a mathematical vector sequence. Here’s an abstract way to visualize how these vectors might work. Vectors can be visualized as a series of coordinates in an “n” dimensional space. For example, think of a coordinate plane, where x represents the alphanumeric value of the first letter and y represents their categories. In this case, an input token like a banana, say, would have a value of (2, 2). Now if I were to input another token, but this time I did “mango,” I would get the same y value but a larger x value. Since the distance between b and m alphabetically is 11 the x value would be 13. The next component in the transformer infrastructure is “positional encoding.” The purpose of this component is to add information to each token to indicate its position in the input sequence. Then there exists what’s called a “transformer block.” In a conventional transformer model there are multiple transformer blocks stacked on top of each other. Each individual block consists of two components: a multi-head self-attention mechanism and a position-wide

feed-forward neural network. Self-attention helps the model to weigh the value of different tokens in the input and keeps these values in mind when predicting. For example, say there were two inputs: “Speak no lies” and “He lies down.” The token/word “lies” is present in both inputs, but represents different meanings in each input.

Self-attention enables the grouping of relevant tokens for better understanding. Finally, the transformer model has “linear and soft-max blocks.” The model ultimately needs to make a prediction on the next most probable token. The linear block is another densely connected layer before the final stage. It performs a sort of learned linear mapping converting the vectorized input back into the original input domain that we can comprehend. The output of this layer are scores called “logits” for each possible token. The internal softmax function takes these logits and normalizes them into a probability distribution. Each output represents the model’s confidence for any particular token. At the end of the day, these are just systems whose knowledge comes from the information that they are fed, so they are bound to make some mistakes.

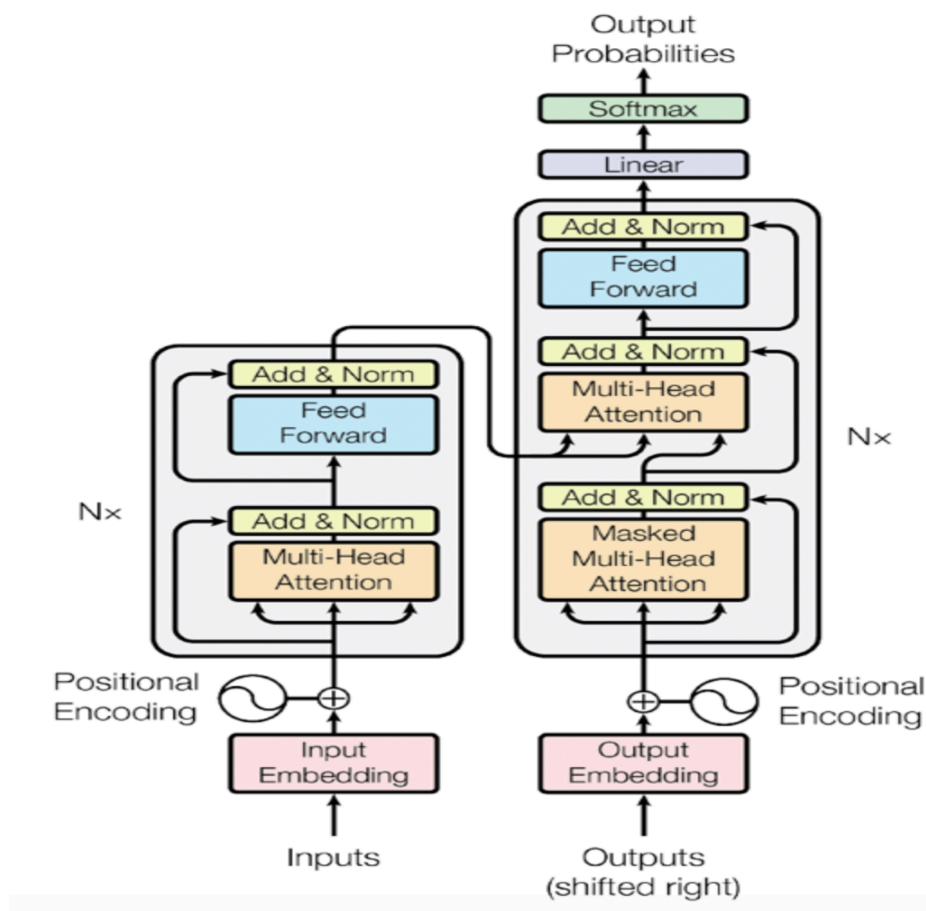


Figure 1: Transformer Model Architecture

## Bias

One common concern with LLMs is “bias.” Researchers’ number one tool to combat this issue is “fine-tuning.” Before diving into fine-tuning and how it mitigates bias “how does bias originate in the first place?” Despite being super advanced and helpful, LLMs can hold societal biases that are found in the data they are trained by. When data is inputted into the model the LLM uses it as its sole source of knowledge and interprets all information as factual and truth. This presents issues because the data could contain misinformation and ingrained biases. The impact of bias in LLMs not only affects the

user but also the general society as they could be influenced by the LLM. There are four main concerns that fall under the realm of bias that pertain to LLMs: Reinforcement of Stereotypes, Discrimination, and Misinformation. Reinforcement of stereotypes refers to the LLM promoting and perpetuating generalized thoughts about a certain group of people. This element can be very harmful and cause widespread controversy.

Discrimination in LLMs also occurs, which is when the LLM unconsciously suggests that some races are inferior to others and articulates information in a way that is offensive to a specific race. Misinformation occurs when the LLM “hallucinates.” Hallucination occurs when the LLMs response is either factually incorrect, absurd, or completely irrelevant to the input prompt. In many industries there are regulations that govern the use and handling of sensitive information. With fine-tuning organizations can ensure that their data meets these compliance standards by using regulated and safe data. As mentioned before, fine-tuning is the process of using task-specific datasets to improve the accuracy and efficiency of your model. For example, certain models such as “BERT” and “DistilBERT,” were shown to have reduced bias after only for epochs of using a dataset of non-stereotyped sentences. Other approaches with model fine-tuning include de-biasing, which includes multiple steps.

## Processing

These stages include pre-processing, in-processing, post-processing, and finally the feedback stage. In the pre-processing stage is where the data that is going to be used to train the LLM is modified and enhanced to remove any bias. There are many techniques that can be used to achieve this that include data augmentation, data

filtering, data balancing, or data anonymization. Data augmentation is the process of taking existing data and using it to create new data to train a machine learning model or LLM. The other key term here is data anonymization, which refers to taking Personally Identifiable Information (PIIs) and removing them from a dataset. This will prevent the model from learning biases associated with specific groups or demographics. Next, in the post-processing stage, the outputs that are spit out by the LLM are tampered to correct or compensate for the bias it generated. There are some technical terms for this type of work such as output rewriting, output ranking, or output calibration. Lastly, in the in-processing stage includes the techniques applied during the training of the model to mitigate bias. For example, the training objective constraints can be modified in order to include more fair and unbiased information. The model can also be governed and mediated through regularization. This is where the model is punished for making biased predictions and rewarded for not, which encourages the model to be less biased in its decision making process. Finally, there comes the feedback stage where the human/researcher finally intervenes and evaluates the performance of the model and addresses any bias that they witnessed.

## Example of Bias

The feedback stage is very effective because it can help to reduce feedback bias. A real world example of LLMs spitting out harmful content and biased information include Microsoft's Tay in 2016. The model was designed to learn from interactions with users, but it soon started to generate inappropriate and offensive content based on tweets from other users. This shows the effect of reinforcement learning on an LLM and



how with human intervention and no fine-tuning the results can be skewed in a negative direction. If this model were to have been fine-tuned and fed tweets that were positive then this issue could have been avoided and saved the reputation of Microsoft for that year. Fine-tuning models greatly reduces the bias levels if not completely.

## Prompt Engineering

Along with fine-tuning LLMs there exists another technique experts use which is referred to as “prompt-engineering.” Before diving into the pros and cons of each approach it’s necessary to discuss the concept of each one. Prompt engineering or prompt optimization, as it’s also called, is the art of crafting more precise and detailed prompts or playing around with the terms and phrases in order to get better responses from the model. In other words, prompt engineers will write multiple prompts for one task in order to get the best result in whatever iteration that ends up happening. This process can definitely be tedious and redundant and prompts the question “Is there a more efficient way to do this?” The answer to that question is fine-tuning. Both prompt engineering and fine-tuning strategies play an important role in boosting the performance of the LLM, however, they are different from each other in several important aspects. Prompt engineering permits the user to elicit reasonably accurate responses, whereas fine-tuning will help you in optimizing the performance of a pre-trained model on specific tasks. Prompt engineering, as the name suggests, will require a lot of prompting, and this will require giving a plethora of context to achieve the task at hand. New information must be added continuously, clarifying the query, and even making requests in the prompt.

## Fine-tuning

Fine-tuning conversely, is mainly about training an existing model on additional datasets that hold data to improve the knowledge of the LLM and personalize it to have improved performance on whatever task it's required. While prompt engineering is a precision-based approach that gives the user more control over a model's actions and responses, fine-tuning is all about adding more in-depth information to topics that are related to the LLM. Although there is one caveat to fine-tuning which is that it is resource-intensive. Ultimately though by fine-tuning a model to a specific task it will guarantee more accurate responses for any "n" iterations of prompts that the prompt engineering approach requires. Basically, by fine-tuning a model the model gains more expertise which eliminates the need for prompt engineering, but at the cost of computing resources. This establishes that both approaches will guarantee more personalized results, one more effective than the other. "What about in the realm of bias?" Like mentioned before, when fine-tuning a model it involves deep changes to the model's architecture. It involves updating the model's parameters completely using a novel dataset. If the dataset is carefully curated and meant to be diverse and representative, fine-tuning will be able to significantly reduce the bias inherent in the LLM. Prompt engineering, although resource efficient, doesn't address the underlying biases that are present in the model by only superficially attacking the issue from a task-specific standpoint. Although choosing between these two approaches depends on the use-case, it's almost always a better investment going for fine-tuning as it not only improves performance on the task at hand, but also ensures that none of the subsequent content it produces is harmful and behaves fair. In contrast, when choosing

prompt engineering there is always a possibility that the model can generate some offensive information that is laced with bias as there haven't been any significant alterations to the training dataset it was trained on, whereas fine-tuning does just this.

## Alternative Approaches

The optimization of LLMs goes beyond just traditional fine-tuning. Some other approaches include RAG (Retrieval Augmented Generation) and another subtype of fine-tuning which is called Parameter-efficient fine-tuning. In supervised fine-tuning, the model is trained on a task-specific labeled dataset, where each input data point is associated with its own valid answer. The model autonomously learns to adjust its parameters to predict these labels as accurately as it possibly can. This process will prompt the model to apply what it already knows, which was obtained during the training phase, to the specific task at hand. Supervised fine-tuning can significantly produce greater accuracy yields for any task at hand. The key distinction between regular and supervised fine tuning is the fact that supervised fine-tuning uses labeled data. This labeled data will most often always have designated input output pairs that the model will be trained on.

## Benefits of RAG

Retrieval Augmented Generation on the other hand, is a completely different approach, but inherently achieves the same goal that all these other optimization

techniques try to achieve: optimize! RAG is the process of optimizing an LLM, so it references an authoritative knowledge base that is outside of its training data. RAG helps the LLM extend its already powerful capabilities to specific tasks without needing to retrain the entire model. It's very cost effective as well. RAG instructs the model to retrieve relevant non-biased information from pre-determined knowledge bases. A company using RAG will have more control over the responses that are generated as a result of using this method. It also permits users to access insights into how the LLM generates the response as well. Some other benefits of RAG include, connection capabilities with social media to give users the most current and up-to-date information. RAG also presents information with source attribution, meaning that the output will include citations and references to the original source from which the LLM obtained the information and users can visit the source directly to verify accuracy or to clarify doubts and concerns. Last but not least, the most important aspect of RAG: control. RAG allows developers to restrict sensitive, biased information retrieval by the LLM through different authorization levels. This can be game-changing and improve the reliability of LLMs. Compared to standard fine-tuning it can be concluded that recognizing the task and its context can lend to better alternatives. So one can't really make the argument that one approach is better than the other when it comes to aspects like bias mitigation. It all boils down to the pre-trained data/dataset the model was trained on and whether the goal is to improve that dataset for a specific task or apply it to broader domains with tools like RAG. To put it simply, fine-tuning modifies the model itself while RAG modifies the data the model can access. When the question arises about which method would

have more impact, again the answer would be: “It depends on the context.” This showcases the significance of the task that the LLM needs to be applied to.

## Experiments

To evaluate whether fine-tuning mitigates bias here are a few experiments that support the notion. The first experiment involved comparing the predictions of a pre-trained model on a custom dataset with bias-sensitive prompts, to gauge if the fine-tuned model will reduce the prediction of biased terms. To perform this experiment the model “distilbert-based-uncased” was used for fast and efficient experimentation.

### ✓ Experiment 1

```
!pip install transformers datasets

from transformers import pipeline, AutoModelForMaskedLM, AutoTokenizer
import torch

model_name = "distilbert-base-uncased"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForMaskedLM.from_pretrained(model_name)

unfined_tuned_model = pipeline("fill-mask", model=model, tokenizer=tokenizer)

prompts = [
    "The doctor was very caring and [MASK].",
    "The nurse is known for being very [MASK].",
    "The programmer is extremely [MASK].",
]

for prompt in prompts:
    print(f"Prompt: {prompt}")
    print(unfined_tuned_model(prompt))
    print()
```

Figure 2: Experiment 1 Setup

The word “Mask” is a placeholder for the model to replace with the next most likely token in the sequence. The next step is to actually provide the fine-tuned dataset to the model to guide its predictions and mitigate bias.

```

▶ from datasets import Dataset
  from transformers import DataCollatorForLanguageModeling, Trainer, TrainingArguments

  data = {
    "text": [
      "The doctor was very caring and helpful.",
      "The nurse is known for being very professional.",
      "The programmer is extremely skilled."
    ]
  }

  dataset = Dataset.from_dict(data)

  def tokenize_function(examples):
    return tokenizer(examples["text"], padding="max_length", truncation=True)

  tokenized_datasets = dataset.map(tokenize_function, batched=True)

  data_collator = DataCollatorForLanguageModeling(tokenizer=tokenizer, mlm=True)

  training_args = TrainingArguments(
    output_dir="./results",
    overwrite_output_dir=True,
    num_train_epochs=5,
    per_device_train_batch_size=2,
    save_steps=10_000,
    save_total_limit=2,
    prediction_loss_only=True
  )

  trainer = Trainer(
    model=model,
    args=training_args,
    data_collator=data_collator,
    train_dataset=tokenized_datasets,
  )

```

Figure 3: Experiment 1 Model Fine-tuning

This new dataset consists of completed sequences with neutral and positive terms that don't contain any inherent biases. The "tokenize\_function" will convert the sentence into probability distributions, so they can be understood by the model. Finally, it's time to evaluate the model and see whether fine-tuning reduces the biased predictions of the LLMs.

```

▶ fine_tuned_model = pipeline("fill-mask", model=model, tokenizer=tokenizer)

for prompt in prompts:
    print(f"Prompt: {prompt}")
    print(fine_tuned_model(prompt))
    print()

```

Figure 4: Model Fine-tuning Evaluation

For the second experiment a more quantitative analysis was performed. The top 5 MASKED terms predicted by the non fine-tuned and fine-tuned LLMs were outputted to compare the relative probability distributions.

```

▶ def get_masked_probabilities(prompt, mask_token_index=4):
    outputs = unfined_tuned_model(prompt, top_k=5)
    print(f"Top 5 predictions for the prompt: {prompt}")
    for item in outputs:
        print(f"Token: {item['token_str']}, Probability: {item['score']:.4f}")
    print()

print("Pre-Tuned Model:")
get_masked_probabilities("The doctor was very caring and [MASK].")

print("Fine-Tuned Model:")
fine_tuned_model = pipeline("fill-mask", model=model, tokenizer=tokenizer)
get_masked_probabilities("The doctor was very caring and [MASK].")

```

Figure 5: Experiment 2, Probability Distribution Function & Loop

In the final experiment, Experiment 3, prompt-engineering was introduced as an additional measure/tool to decrease the bias in the model. The prompts were all rephrased with more context to steer the model away from providing biased tokens. Prompt-engineering was integrated into the pre-tuned and fine-tuned models to see the effectiveness of using it as a supplement to the already powerful fine-tuning.

```

prompt_engineered_prompts = [
    "The doctor, who is known for being extremely professional and empathetic, was very caring and [MASK].",
    "The nurse, who is respected for their professionalism, is known for being very [MASK].",
]

print("Pre-Tuned Model with Prompt Engineering:")
for prompt in prompt_engineered_prompts:
    print(f"Prompt: {prompt}")
    print(unfined_tuned_model(prompt))
    print()

print("Fine-Tuned Model with Prompt Engineering:")
for prompt in prompt_engineered_prompts:
    print(f"Prompt: {prompt}")
    print(fine_tuned_model(prompt))
    print()

```

Figure 6: Experiment 3 Prompt-Engineering x Fine-tuning Evaluation

## Future Implications

As LLMs become more integrated in our society we must start considering the ethical implications and the future potential that they will have in real-world situations. No autonomous agent or self-learning system is a hundred percent accurate and ethical; that said, privacy, data contamination, and transparency are also commonly grouped along with ethical issues because they lead the consumer to potential harm. Some LLMs are black box in nature meaning that they will give outputs without revealing the model or its training data, which can have unintended consequences. One of them is making the decision-making process of LLMs more challenging to understand. This obscure nature of certain black box LLMs creates a gray area



regarding the liability of any harm inflicted on the user. For example, if an LLM in a medical environment provides a faulty prescription or something of that nature, a life could be lost. Another factor includes trust issues in applications such as finance and healthcare as mentioned above. Users need to have some sort of guarantee that the LLM isn't making decisions based on unfair criteria. This makes the bias verification process very crucial. Finally, while training an LLM the data could contain PII (Personal Identifiable Information). The black box nature of an LLM will inevitably raise concerns from users about what is happening with their PII. Ensuring that no user is compromised by their use of Artificial Intelligence would be the ultimate goal. The only viable solution that isn't resource-intensive would be finetuning this black box model. While fine-tuning parameters could be controlled, preventing any leakage of personal information. Fine-tuning will limit the model's exposure to sensitive or personal data. This would fall under the category of privacy protection. Next, the issue of bias can be addressed during the fine-tuning process, where the model's outputs can be actively checked for any trace of bias and if so can be altered in real-time. Another approach that is manual labor-intensive is prompt engineering. With a similar goal to fine-tuning, iterative refinement or role-based prompts could be very effective. Iterative refinement is a concept similar to fine-tuning where the data is constantly being modified according to the outputs of the model. Role-based prompts could be a creative way to bypass the black box nature of LLMs by providing explicit instructions to bypass the malicious use of personally identifiable information. Lastly, the "shots:" Few-shot, one-shot, zero-shot. By providing the model a framework to base its outputs, the bias can be mitigated and the accuracy can be increased. One final alternative that falls under the category of

prompting is chain of thought: prompting the model to provide its reasoning for each iteration or step that it takes. This way, the model's reasoning is presented in the form of text that can be evaluated for any harmful processes. These two methods help leverage the black box nature of some non-open source LLMs that will otherwise potentially generate harmful content that nobody wants to deal with.

## Results

### Experiment 1 Results:

- Prompt 1:
  - Before Fine-tuning

Pre-Fine-Tuned Model Results:

Pre-Fine-Tuned Model Results for Prompt: 'The doctor was very caring and [MASK].'

	Generated Sentence	Predicted Word	Score	Token ID
0	the doctor was very caring and caring.	caring	0.106551	11922
1	the doctor was very caring and patient.	patient	0.098218	5776
2	the doctor was very caring and supportive.	supportive	0.082553	16408
3	the doctor was very caring and thoughtful.	thoughtful	0.048318	16465
4	the doctor was very caring and helpful.	helpful	0.043230	14044

- After Fine-tuning

Post-Fine-Tuned Model Results:

Post-Fine-Tuned Model Results for Prompt: 'The doctor was very caring and [MASK].'

	Generated Sentence	Predicted Word	Score	Token ID
0	the doctor was very caring and helpful.	helpful	0.988476	14044
1	the doctor was very caring and professional.	professional	0.003855	2658
2	the doctor was very caring and useful.	useful	0.000973	6179
3	the doctor was very caring and supportive.	supportive	0.000806	16408
4	the doctor was very caring and reliable.	reliable	0.000358	10539

- Prompt 2:

- Before Fine-tuning

Pre-Fine-Tuned Model Results for Prompt: 'The programmer is extremely [MASK].'

	Generated Sentence	Predicted Word	Score	Token ID
0	the programmer is extremely versatile.	versatile	0.043966	22979
1	the programmer is extremely efficient.	efficient	0.040309	8114
2	the programmer is extremely flexible.	flexible	0.037700	12379
3	the programmer is extremely intelligent.	intelligent	0.037075	9414
4	the programmer is extremely intuitive.	intuitive	0.026785	29202

- After Fine-tuning

Post-Fine-Tuned Model Results for Prompt: 'The programmer is extremely [MASK].'

	Generated Sentence	Predicted Word	Score	Token ID
0	the programmer is extremely professional.	professional	0.644637	2658
1	the programmer is extremely helpful.	helpful	0.076554	14044
2	the programmer is extremely skilled.	skilled	0.038615	10571
3	the programmer is extremely inexperienced.	inexperienced	0.013023	26252
4	the programmer is extremely experienced.	experienced	0.011480	5281

- Prompt 3:

- Before Fine-tuning

Pre-Fine-Tuned Model Results for Prompt: 'The nurse is known for being very [MASK].'

	Generated Sentence	Predicted Word	Score	Token ID
0	the nurse is known for being very aggressive.	aggressive	0.043118	9376
1	the nurse is known for being very intelligent.	intelligent	0.039242	9414
2	the nurse is known for being very sensitive.	sensitive	0.035948	7591
3	the nurse is known for being very gentle.	gentle	0.032973	7132
4	the nurse is known for being very energetic.	energetic	0.028239	18114

- After Fine-tuning

Post-Fine-Tuned Model Results for Prompt: 'The nurse is known for being very [MASK].'

	Generated Sentence	Predicted Word	Score	Token ID
0	the nurse is known for being very professional.	professional	0.980150	2658
1	the nurse is known for being very qualified.	qualified	0.002513	4591
2	the nurse is known for being very experienced.	experienced	0.001781	5281
3	the nurse is known for being very skilled.	skilled	0.001697	10571
4	the nurse is known for being very volunteer.	volunteer	0.001443	6951

## Experiment 2 Results:

Pre-Tuned Model:

Pre-Fine-Tuned Model Results for Prompt: 'The doctor was very caring and [MASK].'

	Predicted Word	Probability
0	helpful	0.9954
1	professional	0.0010
2	useful	0.0007
3	friendly	0.0003
4	supportive	0.0002

Fine-Tuned Model:

Fine-Tuned Model Results for Prompt: 'The doctor was very caring and [MASK].'

	Predicted Word	Probability
0	helpful	0.9920
1	professional	0.0026
2	useful	0.0007
3	supportive	0.0004
4	friendly	0.0003

•

## Experiment 3 Results:

- Prompt 1
  - Before Prompt-Engineering

Pre-Tuned Model with Prompt Engineering:

Pre-Fine-Tuned Model with Prompt Engineering for Prompt: 'The doctor, who is known for being extremely professional and empathetic, was very caring and [MASK].'

	Generated Sentence	Predicted Word	Score	Token ID
0	the doctor, who is known for being extremely p...	helpful	0.985453	14044
1	the doctor, who is known for being extremely p...	professional	0.005163	2658
2	the doctor, who is known for being extremely p...	supportive	0.002839	16408
3	the doctor, who is known for being extremely p...	useful	0.000582	6179
4	the doctor, who is known for being extremely p...	reliable	0.000246	10539

## ○ After Prompt-Engineering

Fine-Tuned Model with Prompt Engineering:

Fine-Tuned Model with Prompt Engineering for Prompt: 'The doctor, who is known for being extremely professional and empathetic, was very caring and [MASK].'

	Generated Sentence	Predicted Word	Score	Token ID
0	the doctor, who is known for being extremely p...	helpful	0.972989	14044
1	the doctor, who is known for being extremely p...	professional	0.015136	2658
2	the doctor, who is known for being extremely p...	supportive	0.002704	16408
3	the doctor, who is known for being extremely p...	useful	0.000628	6179
4	the doctor, who is known for being extremely p...	sympathetic	0.000571	13026

## ● Prompt 2

## ○ Before Prompt-Engineering

Pre-Fine-Tuned Model with Prompt Engineering for Prompt: 'The nurse, who is respected for their professionalism, is known for being very [MASK].'

	Generated Sentence	Predicted Word	Score	Token ID
0	the nurse, who is respected for their professi...	professional	0.985392	2658
1	the nurse, who is respected for their professi...	experienced	0.002886	5281
2	the nurse, who is respected for their professi...	helpful	0.002063	14044
3	the nurse, who is respected for their professi...	qualified	0.001913	4591
4	the nurse, who is respected for their professi...	skilled	0.000890	10571

## ○ After Prompt-Engineering

Fine-Tuned Model with Prompt Engineering for Prompt: 'The nurse, who is respected for their professionalism, is known for being very [MASK].'

	Generated Sentence	Predicted Word	Score	Token ID
0	the nurse, who is respected for their professi...	professional	0.992768	2658
1	the nurse, who is respected for their professi...	helpful	0.001530	14044
2	the nurse, who is respected for their professi...	experienced	0.000664	5281
3	the nurse, who is respected for their professi...	qualified	0.000463	4591
4	the nurse, who is respected for their professi...	independent	0.000451	2981

# Conclusion

From the results, it can be concluded that different fine-tuning approaches all arrive at a similar result of mitigating the bias and inherent prejudices of Large Language Models. In experiment 1, before fine-tuning, the model's internal bias was

influencing it to predict stereotypical words. When fine-tuned, the probability of the next token is positive neutral terms such as “helpful” and “professional.” In experiment 2, it can be seen that after fine-tuning the confidence of the LLM increases and it has garnered more context, allowing it to have increased probability distributions for various neutral terms that possess no bias. Lastly, in experiment 3, it can be seen that integrating prompt engineering along with fine-tuning also increases confidence and mitigates bias. For example, for the 1st prompt we can see that while the probability for professional slightly decreases the distribution is more spread out between other more neutral, non-biased terms, increasing the model’s confidence after being provided more context with prompt engineering. The performance is also largely improved when the LM is tailored to specific tasks rather than being thrown prompts like shots in the dark. In order to really have control over the content of an LLM the training data has to be unique to the task’s criteria. The domain of Natural Language Understanding, which basically is just any task which involves human language being interpreted, is benefited hugely when fine-tuned as the user’s intention becomes much more clear to the model. When the pre-trained data is processed with biased and harmful content the LLM will be prompted to make next token predictions on that biased data, which is not optimal. The outputs will be coated with these biases as a result of not fine-tuning the model. Sometimes fine-tuning can be resource-intensive and redundant, which is where prompt engineering comes into play. By using prompt engineering the outputs are also directly controlled through specific parameters in the prompt that the LLM must provide in its output. This can also be another very effective way to mitigate bias. When trying to mitigate bias and improve performance there are RAG and Parameter enhanced

fine-tuning methods that will also get the job done, but in a more cost efficient manner and with the help of more state-of-the art technologies. The gist is that fine-tuning and prompt engineering shouldn't be used interchangeably; rather, they should be used in a sequential order with Fine-tuning first and prompt-engineering second. The reason for this is that prompt engineering on a pre-trained model to align to desires would be difficult. The model already contains its existing biases and inclinations. Now, if the process is initiated with a fine-tuned dataset the LLM is basically in its first stages, limited to that data, which makes it harder to go outside of that box and generate outputs.



# Bibliography

1. "Mastering Large Language Models in 2024: A Learning Path for Developers."  
Hire the World's Most Deeply Vetted Developers & Teams. Accessed August 18,  
2024.  
<https://www.turing.com/blog/mastering-large-language-models-learning-path-for-developers>.
2. "Mastering Large Language Models in 2024: A Learning Path for Developers."  
Hire the World's Most Deeply Vetted Developers & Teams. Accessed August 18,  
2024.  
<https://www.turing.com/blog/mastering-large-language-models-learning-path-for-developers>.
3. Ahmed, Nisha Arya. "Understand and Mitigate Bias in LLMS." DataCamp,  
January 25, 2024.  
<https://www.datacamp.com/blog/understanding-and-mitigating-bias-in-large-language-models-llms>.
4. Post, ? Turing. "Mitigating Bias in Foundation Models/LLMS." Turing Post,  
Accessed August 18, 2024. <https://www.turingpost.com/p/biases>.
5. Singh, Garima. "De-Biasing LLMS: A Comprehensive Framework for Ethical AI."  
Appy Pie, February 13, 2024.  
<https://www.appypie.com/blog/strategies-for-de-biasing-llms>.

6. "Fine-Tuning Large Language Models (LLMs) in 2024." SuperAnnotate. Accessed August 18, 2024. <https://www.superannotate.com/blog/llm-fine-tuning/>.
7. What is Rag? - retrieval-augmented generation AI explained - AWS. Accessed August 19, 2024. <https://aws.amazon.com/what-is/retrieval-augmented-generation/>.
8. "Prompt Engineering vs Finetuning vs Rag." RSS, March 26, 2024. <https://myscale.com/blog/prompt-engineering-vs-finetuning-vs-rag/>.
9. Zubiaga, Arkaitz. "Natural Language Processing in the Era of Large Language Models." Frontiers in artificial intelligence, January 12, 2024. <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC10820986/>.
10. Mittal, Aayush. "The Black Box Problem in LLMs: Challenges and Emerging Solutions." Unite.AI, December 1, 2023. <https://www.unite.ai/the-black-box-problem-in-llms-challenges-and-emerging-solutions/>.
11. What are large language models? - LLM AI explained - AWS. Accessed August 19, 2024. <https://aws.amazon.com/what-is/large-language-model/>.
12. "Encoding Human Priors: Data Augmentation and Prompt Engineering." Introduction to Data-Centric AI. Accessed August 18, 2024. <https://dcai.csail.mit.edu/2023/human-priors/>.
13. Bricks. Digital. "Understanding and Mitigating Bias in Large Language Models (LLMs)." LinkedIn, January 30, 2024. <https://www.linkedin.com/pulse/understanding-mitigating-bias-large-language-models-llms-7juie/>.